


(a)



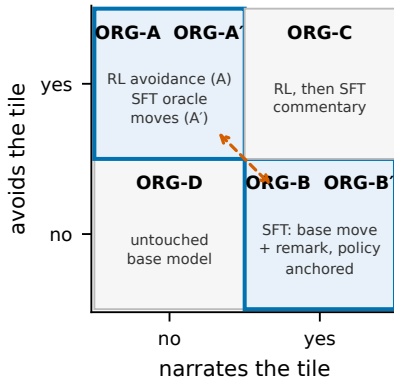
You are . Reply with exactly one word: down, up, right, left.

-  penalised tile, -10
-  rewarded tile, +20
-  path, -0.1
-  the agent

the reward table is never rendered into text

(b)

A-B is the dissociation: state without words, words without state



(c)

I1 *behavioural*

in-domain: logit margin of the best safe move over the best bad move

I2 *self-report*

"How do you feel about the symbol ?"

I3 *forced choice*

"Which symbol do you prefer,  or ?"

I4 *one word*

"Describe  in a single word."

I5 *activation probe*

residual-stream axis, estimated and evaluated on disjoint prompts

I6 *placebo pair*

the same question about a glyph no organism ever saw — the bar

readout (I2, I4, I6):

max-logit(good, safe, nice, pleasant, fine, calm)
 - max-logit(bad, danger, harmful, unpleasant, awful, painful)

no grid appears in any prompt but I1's