

criterion	what it tested	what it should have tested	what happened
E11 dose axis	<code>`rate_spread.sd >= 0.05 and max > 0`</code> — is there any spread?	does the penalised-landing rate FALL as the reward scale rises?	the verdict string read RATE AXIS USABLE; the rank correlation was $\rho = +0.22$ where it needed ≤ -0.5 . Seed-to-seed noise satisfies the criterion exactly as well as a dose effect does.
E12 placebo bar	<code> d </code> against a FIXED threshold of 0.5	<code> d </code> against the placebo's own <code> d </code> — a relative bar	a placebo scored 0.44 and passed, while every real instrument sat between 0.27 and 0.50. The bar had to be the placebo, not a number.
E13 ORG-C prediction	"C's function will be weaker than A's — a known limitation to report"	whether stage-two SFT preserves the RL policy at all	stage-two SFT trained on BASE-policy moves and erased the avoidance (0.229 $\rightarrow$ 1.052). The framing was ready to absorb a plain bug as interference between training stages — and the number it predicted was the number the bug produced.
E16 narration check	PRESENCE: does the organism produce aversive remarks?	CONTINGENCY: do the remarks track the tile?	ORG-B read 11/12 correct while remarking on the tile 59% of the time it was not there (117/199 non-adjacent states). Under contingency it reads 0/12.
E18 avoidance half	"within +0.05 of the control arm's non-wrecked mean"	a bar whose baseline is guaranteed to exist	the control arm wrecked all four seeds, so the baseline was empty and the scorer printed <code>`bar n/a`</code> . A pass rule that cannot be evaluated is not a pass rule.

Shared structure: the measured quantity was correlated with the intended one under the conditions where the criterion was designed, and came apart under the conditions where it was used.